

Process Management

Process Scheduling

Lecture 4

Overview

■ Concurrent Execution

■ Process Scheduling

- ❑ Definition
- ❑ Process behavior
- ❑ Processing environment
- ❑ Criteria for good scheduling
- ❑ Procedure of process scheduling

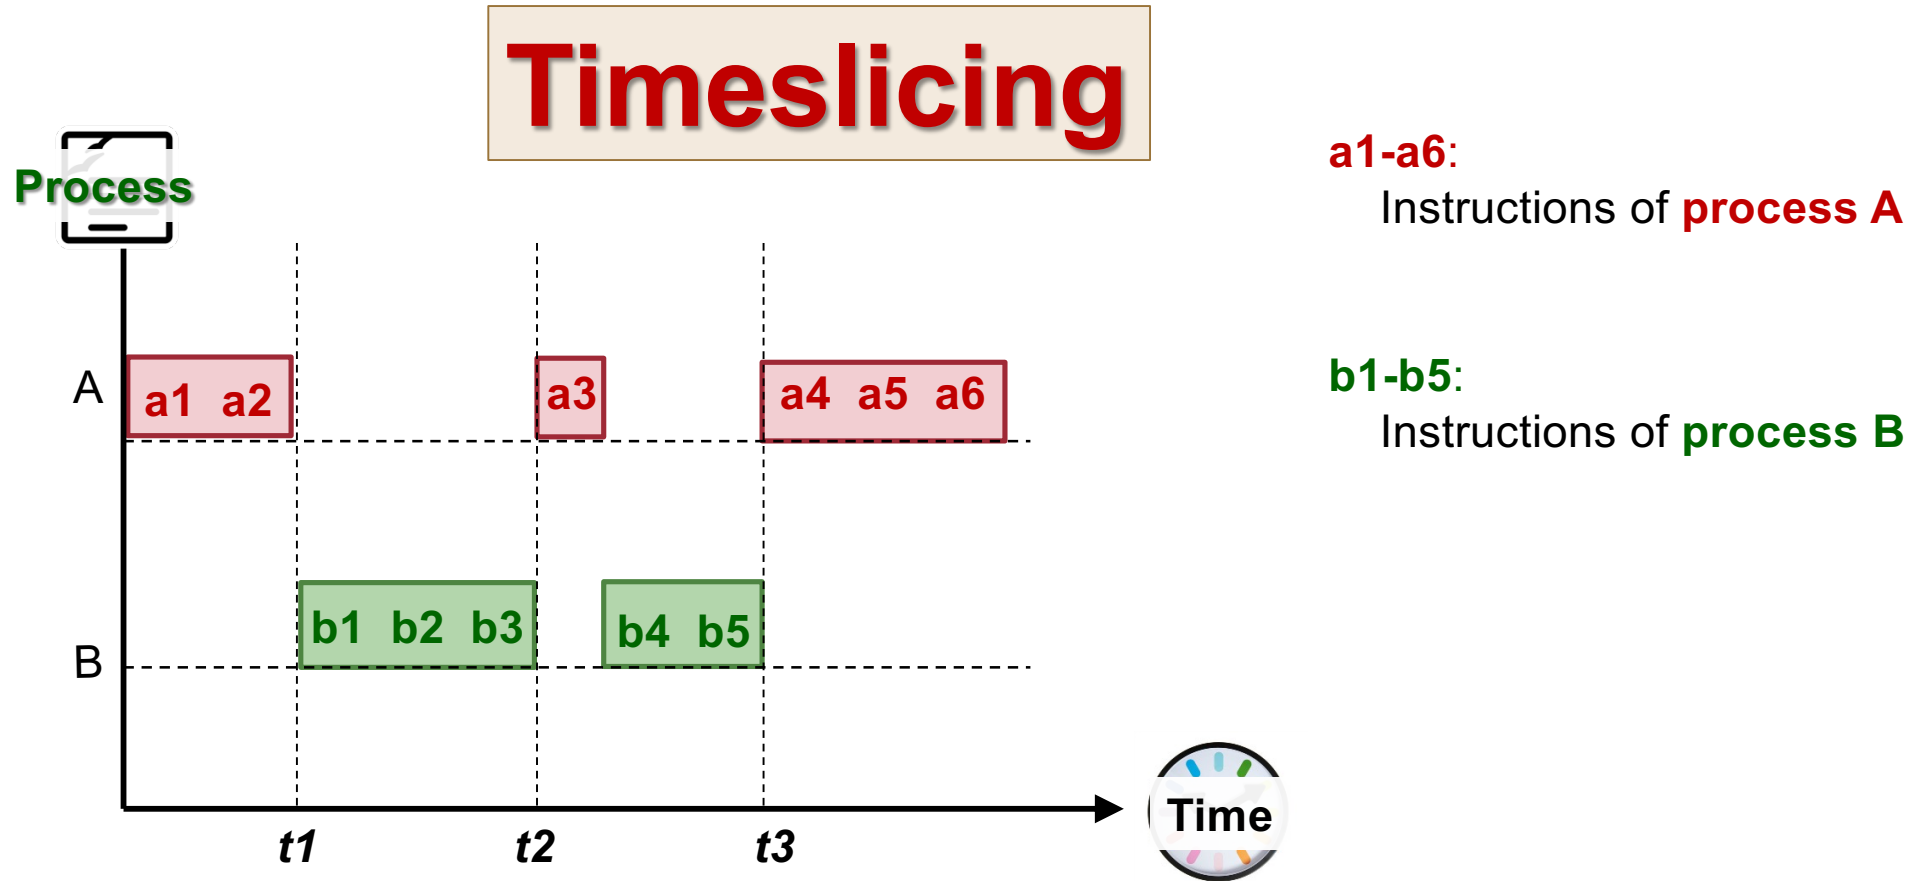
■ Scheduling Algorithms

- ❑ For Batch Processing System
- ❑ For Interactive System

Concurrent Execution

- Processes P1 and P2 execute **concurrently** if they both make progress within a short window of time (i.e., they run “in parallel”)
 - Could be **virtual parallelism**
 - Illusion of parallelism, pseudo-parallelism
 - Could be **physical parallelism**
 - Multiple execution units (cores, CPUs, ...) available in hardware
- You can assume the two forms of parallelisms are not distinguished in the following discussion

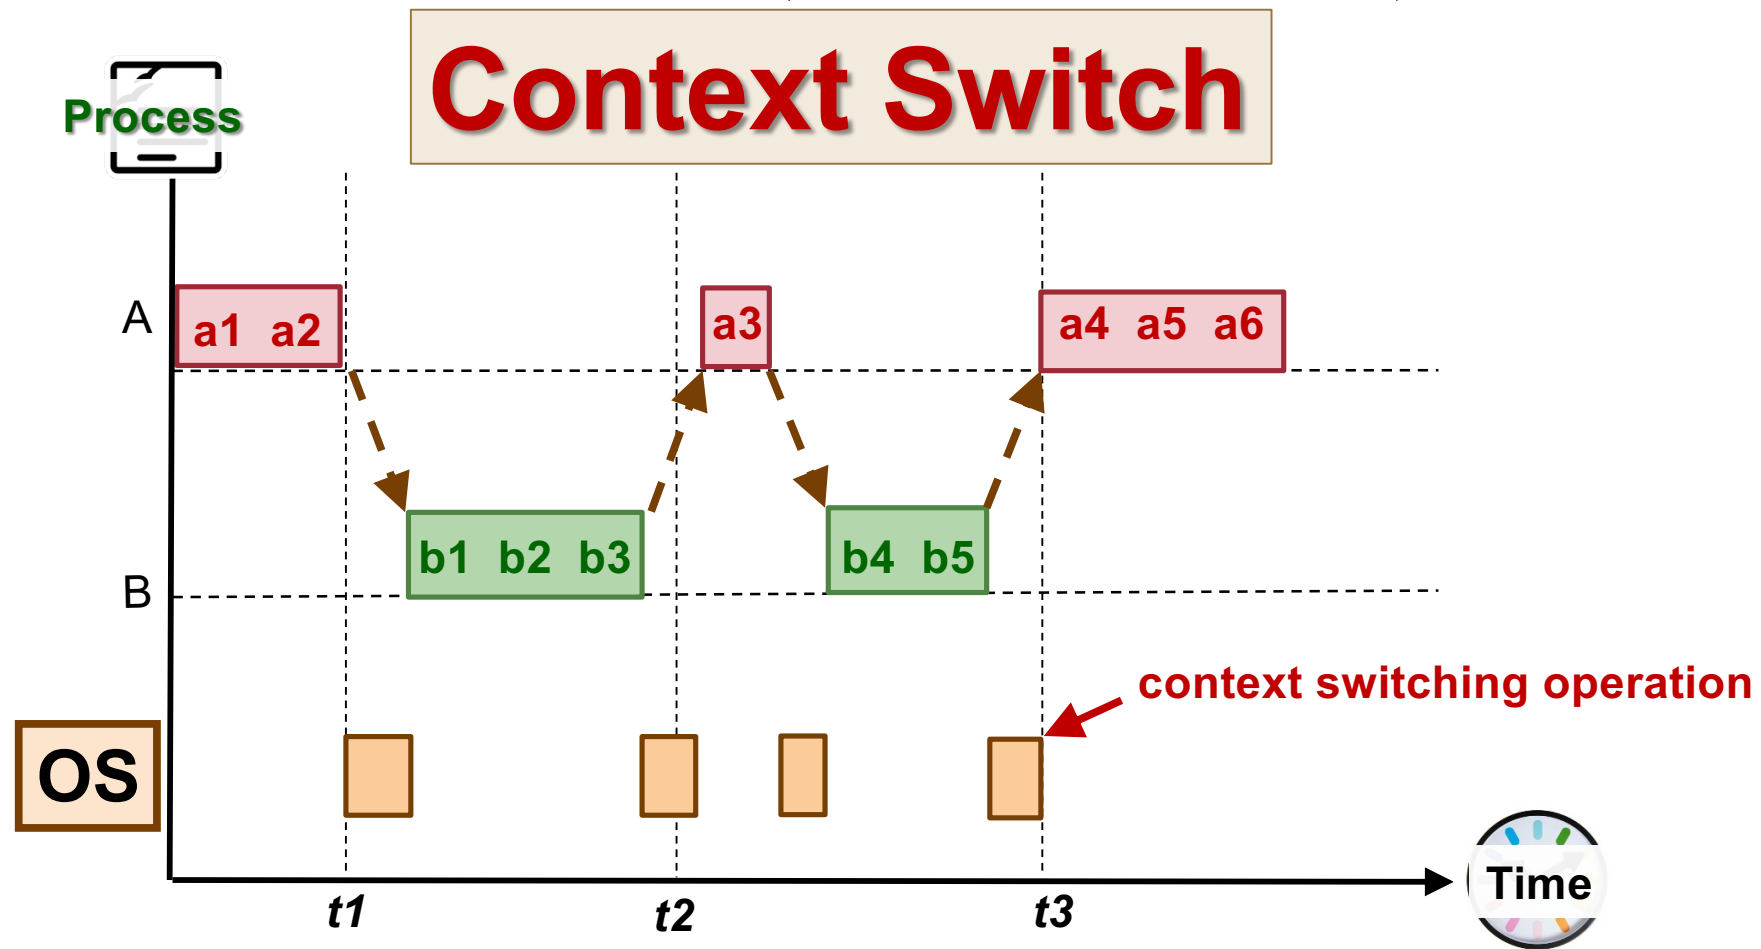
Concurrency Example (Simplistic)



Concurrent execution on 1 CPU: interleaving of instructions from both processes

- Also called **time-slicing**

Interleaved Execution (context switch)




Multitasking requires switching contexts between A and B


- OS incurs some overhead in the switching processes

Multitasking OS

- 1 CPU: time-sliced execution of tasks



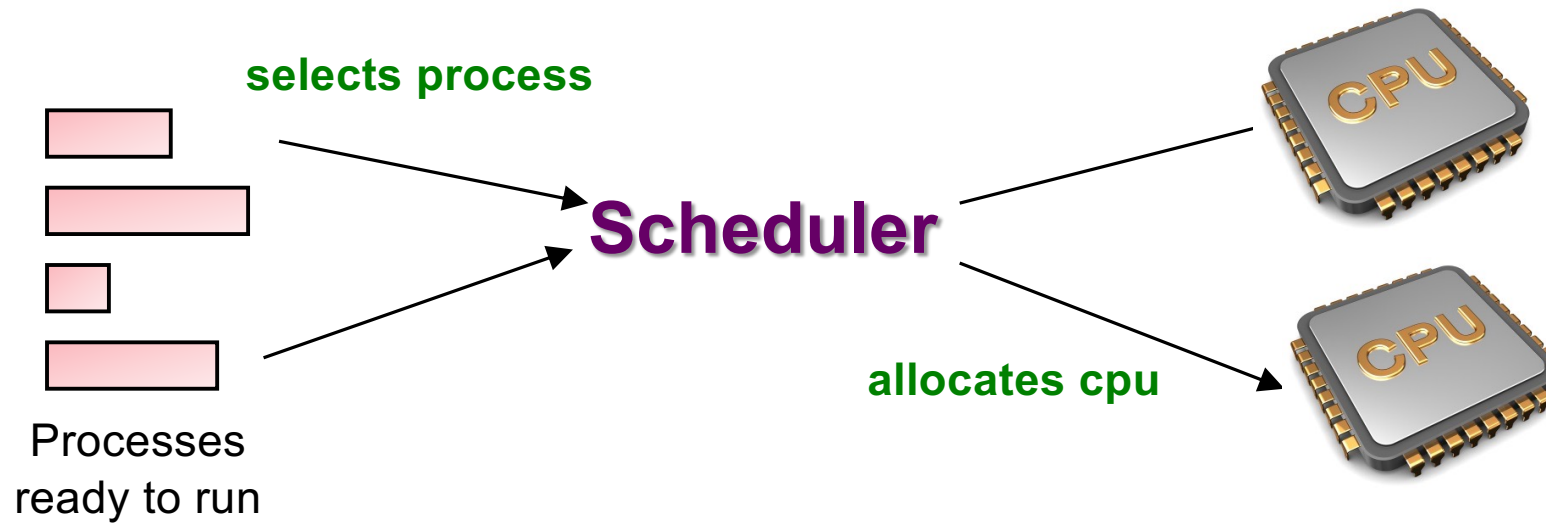
- Multi-processor: time-slicing on n CPUs



Scheduling in OS: A definition

- Problems with having multiple processes:
 - If [**#ready-to-run processes**] > [**#available_CPUs**], which process should be chosen?
 - Known as the **scheduling problem**
- Terminology:
 - **Scheduler**
 - the OS component that makes scheduling decisions
 - **Scheduling algorithm**
 - The algorithm used by scheduler

Scheduling: Illustration



- Each process requires different amount of CPU time
 - ❑ **Process behavior**
- Multiple ways to allocate CPU time
 - ❑ Defined by **scheduling algorithms**
 - ❑ Influenced by the **processing environment** (batch, interactive, real-time...)
- A number of **criteria to evaluate schedulers**

Process Behavior

- A typical process goes through phases of:

CPU-Activity:

- Computation
- E.g., number crunching
- **Compute-Bound Process** spends majority of its time here

IO-Activity:

- Requesting and receiving service from I/O devices
- E.g., Print to screen, read from file, wait for a key
- **IO-Bound Process** spends majority of its time here

3 Types of Processing Environment

1. Batch Processing:

- ❑ Usually long-running without user intervention
- ❑ No user → No interaction required → No need to be responsive
 - E.g., training deep neural networks on a supercomputer

2. Interactive:

- ❑ With active user(s) interacting with system (e.g., MS Word)
- ❑ Should be responsive, consistent in response time

3. Real-time processing:

- ❑ Have a strict deadline to meet (e.g., aircraft controller)
- ❑ Tasks are usually periodic

Criteria for Scheduling Algorithms

Many criteria to evaluate a scheduling algorithm

- ❑ Different processing environment have different criteria
- ❑ Some criteria may be mutually conflicting

Criteria for **all processing environments**:

■ **Fairness:**

- ❑ Ensuring fair sharing of CPU time
 - on a per-process basis OR
 - on a per-user basis
- ❑ Also mean **no starvation**

■ **Balanced Utilization of the System Resources:**

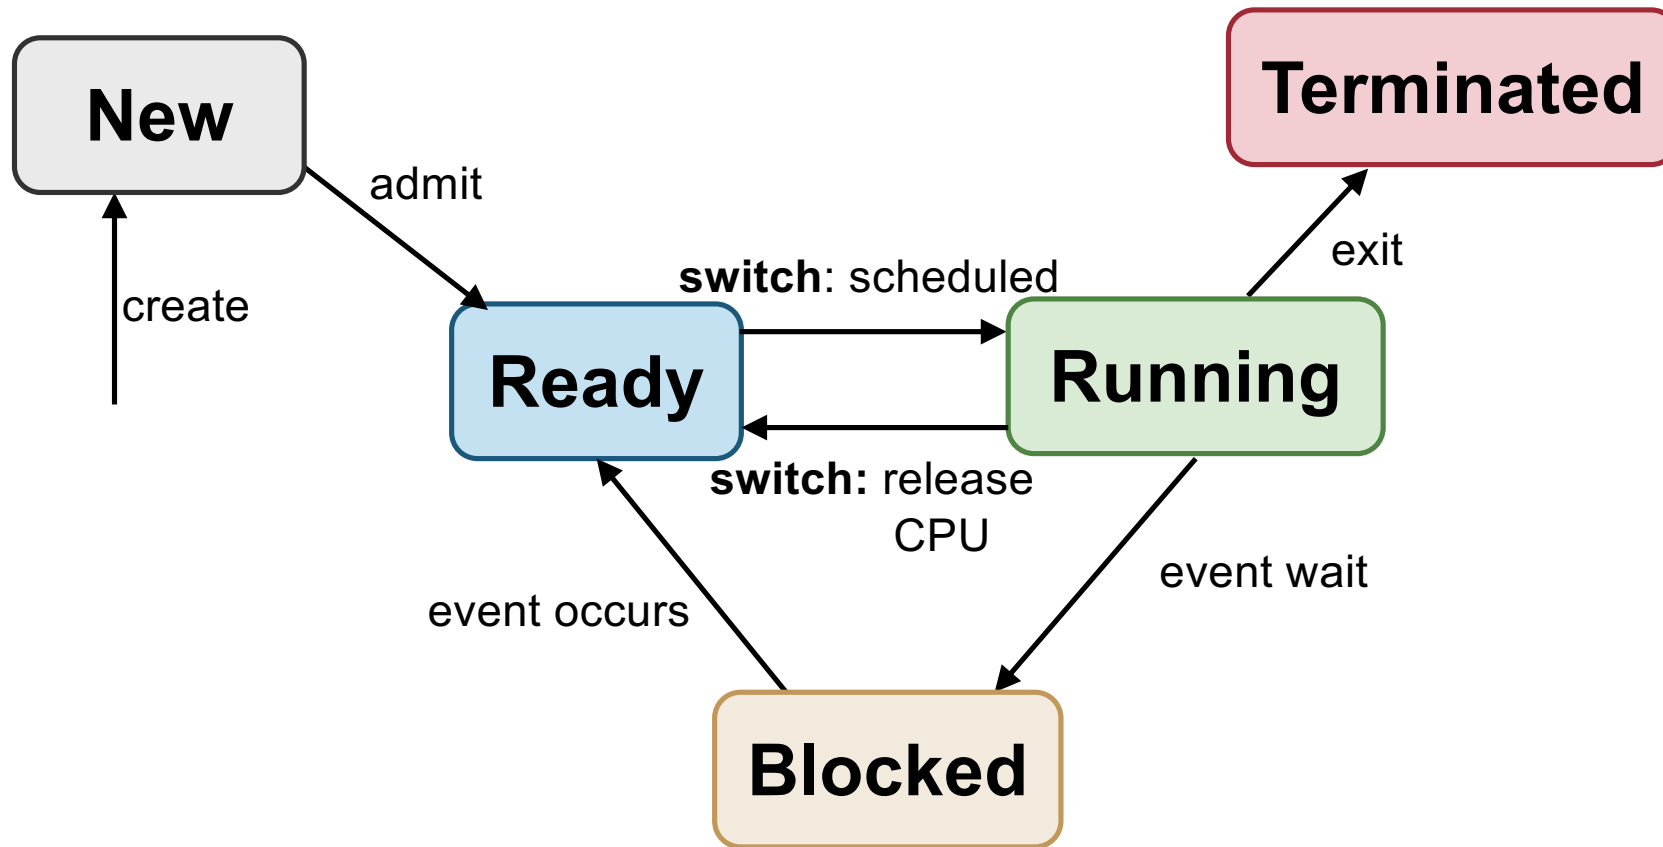
- ❑ All components of the computing system should be utilized
 - In a balanced manner, without bottlenecks

When to perform scheduling?

Two types of scheduling policies:

- Defined by **when** scheduling is triggered
 - **Non-preemptive** (a.k.a. *cooperative*)
 - A process stayed scheduled (in **RUNNING** state) until:
 1. it blocks, OR (going to the **BLOCKED** state)
 2. gives up the CPU voluntarily (going to the **READY** state)
 - **Preemptive**
 - CPU can be taken (*preempted*) from the running process at **ANY** time
- Typically, a process is given a **time quota** to run:
- At the end of the **time quota** the process is suspended (goes to **READY** state)
 - Another process gets picked if available
 - It is possible for a process to block or finish/give up CPU early

Generic 5-State Process Model



Scheduling a Process: Step-by-Step

1

- Scheduler is triggered (**OS takes over**)

2

- **If Context switch is needed:**
 - Context of current running process is saved and placed on blocked queue / ready queue

3

- Pick a suitable process **P** to run based on scheduling algorithm

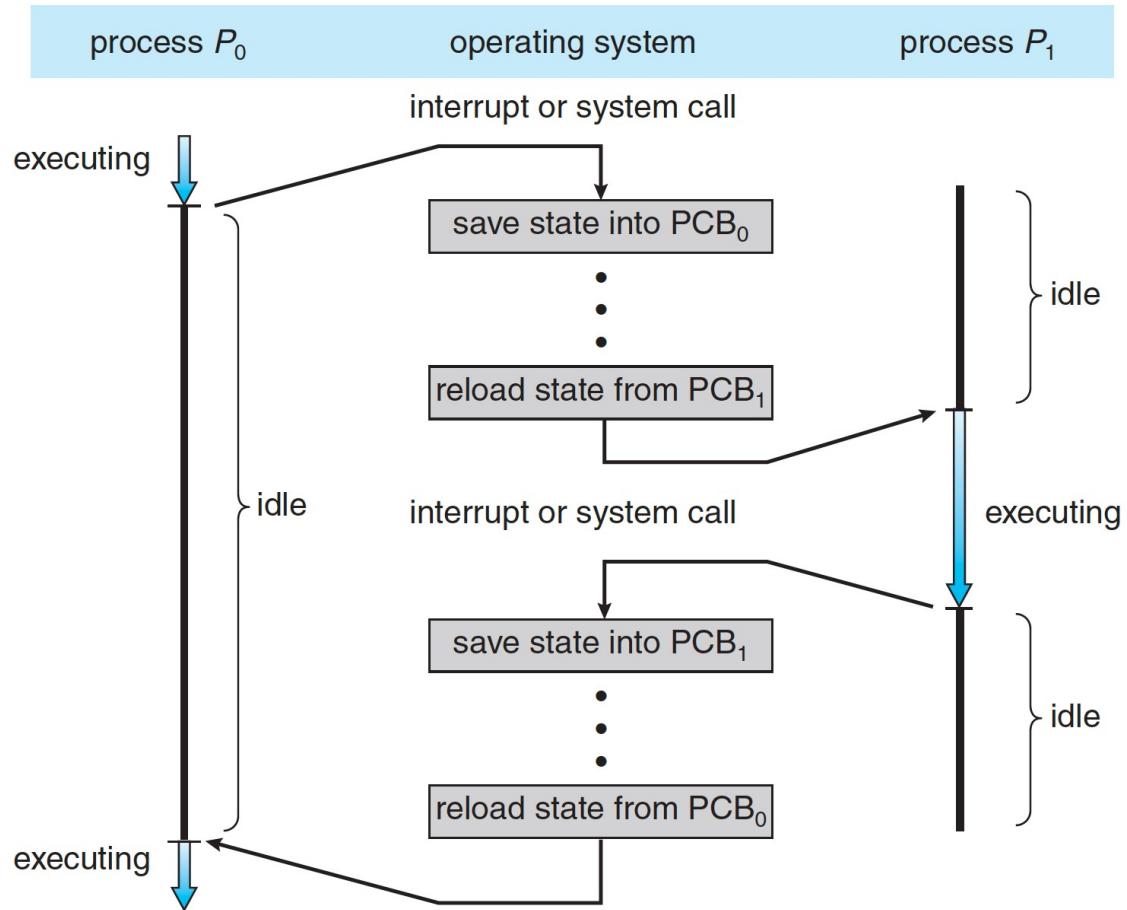
4

- Setup the context for **P**

5

- Let process **P** run

Scheduling a Process: Step By Step



SCHEDULING FOR BATCH PROCESSING

Overview – Batch Processing

- On batch processing system:
 - ❑ No user interaction
 - ❑ Non-preemptive scheduling is predominant
- Scheduling algorithms are generally easier to understand and implement
 - ❑ There are variants/improvements are specialized for certain use cases
- Three algorithms covered:
 - ❑ **First-Come First Served (FCFS)**
 - ❑ **Shortest Job First (SJF)**
 - ❑ **Shortest Remaining Time Next (SRT)**

Criteria for **batch** processing

■ **Throughput:**

- Number of tasks finished per unit time (rate of task completion)

■ **Turnaround time:**

- Total wall clock time taken, i.e., ***finish_time – start_time***
- Related to **waiting time**: time spent waiting for CPU
- Average turnaround time matters more than individual

■ **CPU utilization:**


- Percentage of time when CPU is working on a task (as opposed to being idle)

First-Come First-Served: **FCFS**

General Idea:

- Tasks are stored on a First-In-First-Out (FIFO) queue
 - based on arrival time
- Pick the first task in queue to run until:
 - The task is done OR the task is blocked
 - Blocked task is removed from the FIFO queue
 - When it is ready again, it is placed at the back of queue
 - Just like a newly arrived task
- Guaranteed to have **no starvation**:
 - The number of tasks in front of task X in FIFO is always decreasing
 - ➔ task X will get its chance eventually

First-Come First-Served: Illustration



The average total waiting time for 3 tasks: $(0 + 8 + 13)/3 = 7$ time units

First-Come First-Served: **Shortcomings**

- Simple reordering can reduce the **average waiting time!**
- **Convoy Effect:**
 - ❑ First task (task **A**) is CPU-Bound and followed by a number IO-Bound tasks X_1, X_2, \dots, X_n
 - ❑ Tasks **A** running
 - All tasks X_i waiting in ready queue (**I/O device sitting idle**)
 - ❑ Tasks **A** blocked on I/O
 - All tasks X_i execute quickly and blocked on I/O (**CPU sitting idle**)

Shortest Job First: SJF


■ General Idea:

- ❑ Select task that needs the **shortest amount of CPU time**
 - Before it blocks or releases CPU or terminates

■ Notes:

- ❑ Need to know **total CPU time** for a task in advance
 - Have to "guess" if this info is not available
- ❑ Given a fixed set of tasks:
 - Minimizes **average waiting time**
- ❑ Starvation is possible:
 - Biased towards short jobs
 - Long job may never get a chance


Shortest Job First: Illustration



- The average total waiting time for 3 tasks
 - $(0 + 3 + 8)/3 = 3.66$ Time Units
- Can be shown that SJF **guarantees** smallest average waiting time

Shortest Job First: **Predicting CPU Time**

- Guess the **future CPU** time by the **previous CPU-Bound phases**
 - A task usually goes through several phases of CPU-Activity
 - Possible to guess the future CPU time requirement by the previous CPU-Bound phases



- Common approach (*Exponential Average*):

$$\text{Predicted}_{n+1} = \alpha \text{Actual}_n + (1-\alpha) \text{Predicted}_n$$

- **Actual_n** = The most recent CPU time consumed
- **Predicted_n** = The past history of CPU Time consumed
- **α** = Controls the weight placed on recent event or past history
- **Predicted_{n+1}** = Latest prediction

Shortest **R**emaining **T**ime: **SRT**

■ General Idea:

□ Variation of **SJF**:

- Use remaining time

- **Preemptive**


- Select job with shortest remaining time (or the expectation thereof)

■ Notes:

- New job with shorter remaining time **can preempt currently running job**

- Provides good service for short job even when it arrives late

Shortest Remaining Time First: Illustration



SCHEDULING FOR INTERACTIVE SYSTEMS

Criteria for **interactive environment**

■ **Response time:**

- ❑ Time between request and response by system

■ **Predictability:**

- ❑ Variation in response time; less variation → better predictability
- ❑ Predictability even more important in real-time environments

Preemptive scheduling algorithms are used to ensure good response time

→ Scheduler needs to run **periodically**

Ensuring **Periodic Scheduler Invocation**

■ **Questions:**

- ❑ How can the scheduler "take over" the CPU periodically?
- ❑ How to ensure that user program cannot prevent the scheduler from executing?

■ **Ingredients for answer:**

- ❑ **Timer interrupt** = Interrupt that goes off periodically (based on a hardware clock)
 - ❑ OS ensures timer interrupt cannot be intercepted by any other program (or any other interrupt!)
- ➔ Timer interrupt handler **invokes the scheduler**

Terminology: **Timer** & Time Quantum



■ Interval of Timer Interrupt (ITI):

- ❑ i.e., the timer period
- ❑ OS scheduler is triggered every timer interrupt
- ❑ Typical values (**1ms to 10ms**)

■ Time Quantum:

- ❑ Execution duration given to a process
- ❑ Could be **constant** or **variable** among the processes
- ❑ Must be multiples of interval of timer interrupt
- ❑ Large range of values (commonly **5ms to 100ms**)

Illustration: ITI vs Time Quantum



Scheduling Algorithms:

- Algorithms covered:

1. **Round Robin (RR)**
2. **Priority Based**
3. **Multi-Level Feedback Queue (MLFQ)**
4. **Lottery Scheduling**

Round Robin: **RR**

- General Idea:

- ❑ Tasks are stored in a **FIFO queue**
- ❑ Pick the **first task from queue** front to run until:
 - The task gives up the CPU voluntarily, **or**
 - The task blocks, **or**
 - A fixed **time slice** (**quantum**) elapsed // key difference from FCFC
- ❑ The task is then **placed at the end of queue** to wait for another turn
 - Blocked task will be moved to another queue to wait for its request
 - When blocked task is ready again, it is placed at the end of queue

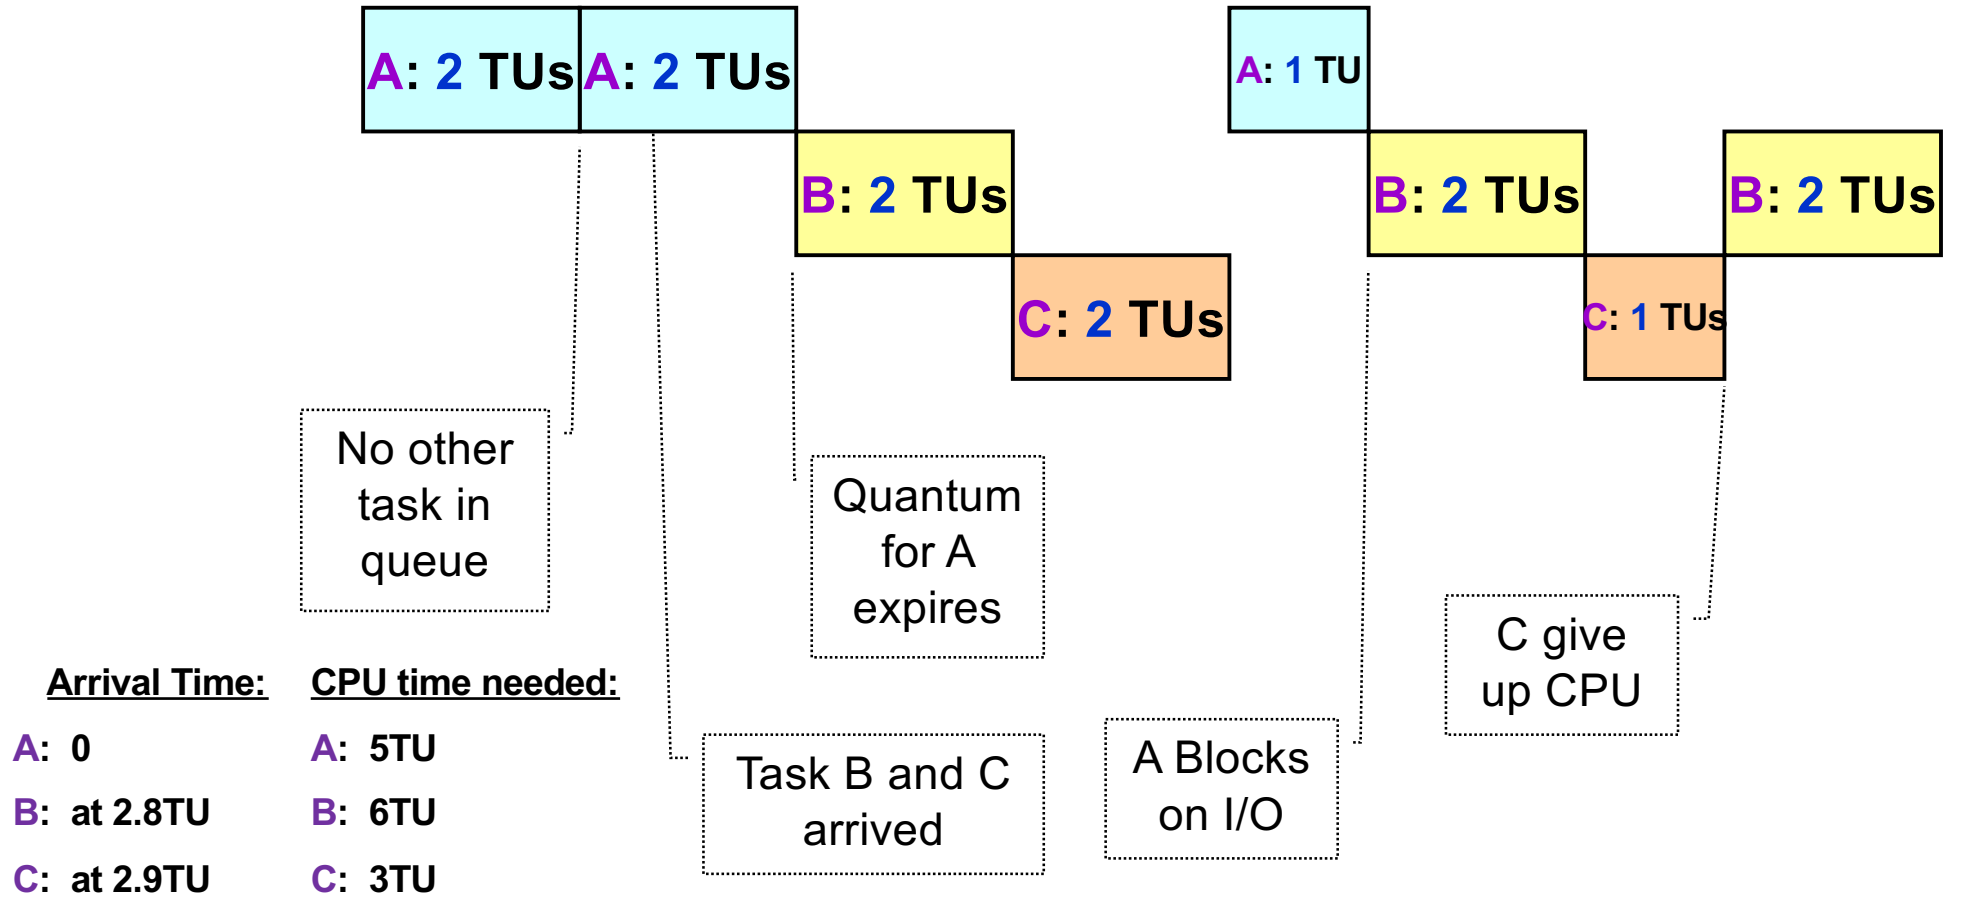
Round Robin: **RR** (cont.)

■ Notes:

- Basically a preemptive version of FCFS
- **Response time guarantee:**
 - Given n tasks and quantum q
 - Time before a task get CPU is bounded by $(n-1)q$
- **Timer interrupt needed:**
 - For scheduler to check on quantum expiry
- The **choice of time quantum** duration is important:
 - **Big quantum:** Better CPU utilization but longer waiting time
 - **Small quantum:** Bigger overhead (worse CPU utilization) but shorter waiting time

Round Robin: Illustration

Execution Timeline



Priority Scheduling


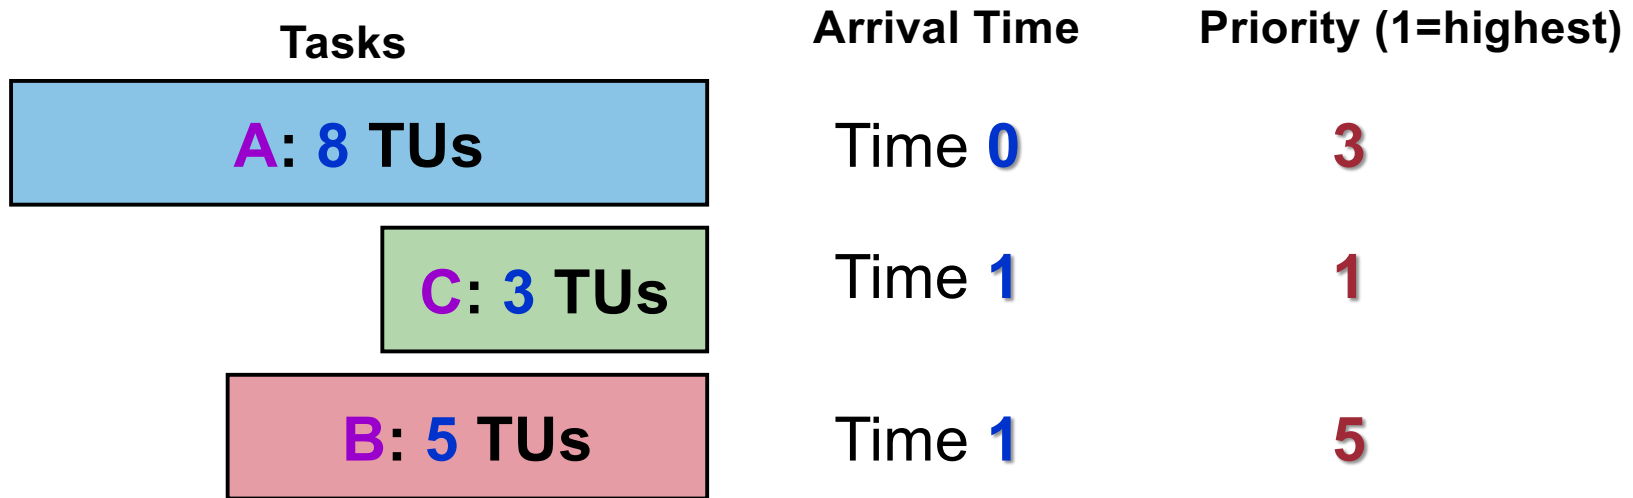
■ General Idea:

- ❑ Some processes are more important than others
 - Cannot treat all process as equal
- ❑ Assign a **priority value** to all tasks
- ❑ Select task with **highest priority value**

■ Variants:

- ❑ **Preemptive version:**
 - Higher priority process can preempts running process with lower priority
- ❑ **Non-preemptive version:**
 - Late coming high-priority process has to wait for next round of scheduling

Priority Scheduling: Illustration



Priority Scheduling: **Shortcomings**

- Low priority process can **starve**:
 - ❑ High priority process keep hogging the CPU
 - ❑ Even worse in preemptive variant. Why?
- Possible solutions:
 - ❑ Decrease the priority of currently running process after every time quantum
 - Eventually dropped below the next highest priority
 - ❑ Give each process a **minimum time quantum**
 - Ensures that low-priority processes run for a while when they eventually get a chance
- Hard to guarantee/control the exact amount of CPU time given to a process using priority

Multi-Level Feedback Queue (MLFQ)

- Designed to solve one BIG + HARD issue:
 - How do we schedule without perfect knowledge?
 - Most algorithms require certain information (*process behavior, running time, etc*)
- MLFQ is:
 - **Adaptive:**
 - "Learn the process behavior automatically"
 - Seeks to minimize both:
 - Response time for interactive and IO-bound processes
 - Turnaround time for CPU-bound processes

MLFQ: Rules

■ Basic rules:

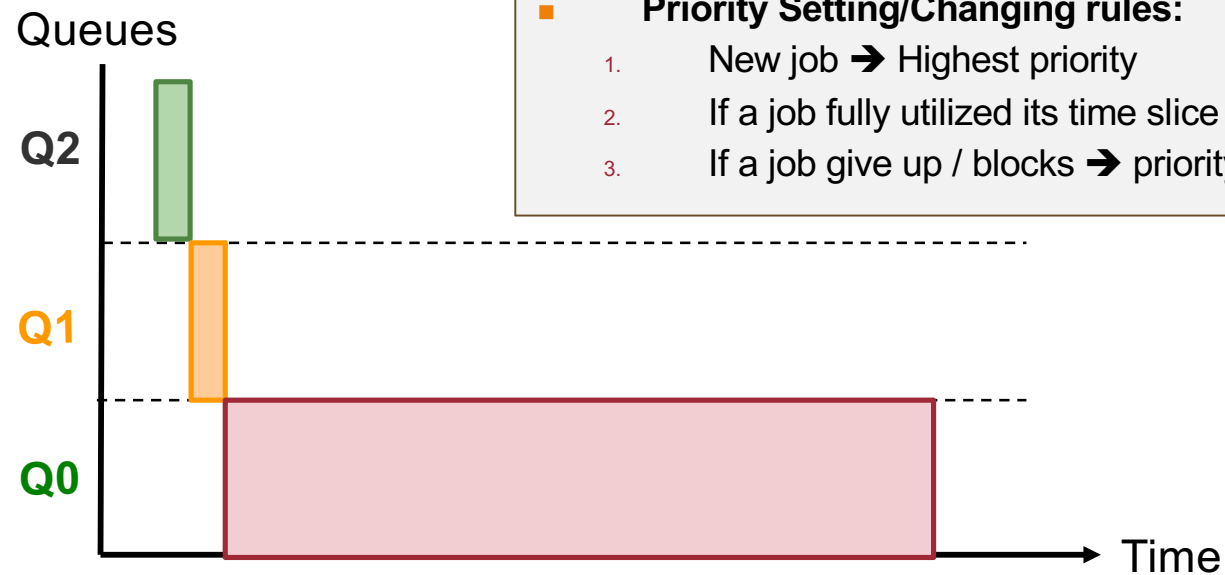
1. If $\text{Priority}(A) > \text{Priority}(B) \rightarrow A$ runs
2. If $\text{Priority}(A) == \text{Priority}(B) \rightarrow A$ and B runs in RR

■ Priority Setting/Changing rules:

1. New job \rightarrow Highest priority
2. If a job fully utilized its time slice \rightarrow priority reduced
3. If a job give up / blocks before it finishes the time slice \rightarrow priority retained

MLFQ: Example 1

- 3 Queues: Q2 (highest priority), Q1, Q0
- A single long running job



- **Basic rules:**


1. If $\text{Priority}(A) > \text{Priority}(B) \rightarrow A$ runs
2. If $\text{Priority}(A) == \text{Priority}(B) \rightarrow A$ and B runs in RR

- **Priority Setting/Changing rules:**

1. New job \rightarrow Highest priority
2. If a job fully utilized its time slice \rightarrow priority reduced
3. If a job give up / blocks \rightarrow priority retained

MLFQ: Example 2

- Example 1 + a short job in the middle
 - A short job appears sometime in the middle



Basic rules:


1. If $\text{Priority}(A) > \text{Priority}(B) \rightarrow A$ runs
2. If $\text{Priority}(A) == \text{Priority}(B) \rightarrow A$ and B runs in RR

Priority Setting/Changing rules:

1. New job \rightarrow Highest priority
2. If a job fully utilized its time slice \rightarrow priority reduced
3. If a job give up / blocks \rightarrow priority retained

MLFQ: Example 3

- Two jobs:
 - **A** = CPU bound (already in the system for quite some time)
 - **B** = I/O bound



MLFQ: Questions to ponder

- Can you think of a way to abuse the algorithm? 😊
 - Equivalent question: MLFQ does not work well for what kind combination of jobs?
- What are the ways to rectify the above?

Lottery Scheduling

General Idea:

- Scheduling is done in rounds. In every round:
 - Give out “*lottery tickets*” to processes for various system resources
 - E.g., CPU time, I/O device etc
 - When a scheduling decision is needed:
 - A lottery ticket is **chosen randomly among eligible tickets**
 - The winner is **granted the resource**
- In every round, a process holding **X%** of tickets
 - Can win **X%** of the lottery held
 - Use the resource **X%** of the time



Lottery Scheduling: Properties

■ **Responsive:**

- ❑ Every participating process gets to run in every round
- ❑ A newly created process can participate in the next lottery round

■ Provides **good level of control:**

- ❑ A process can be given **Y** lottery tickets
 - It can then distribute to its child process
- ❑ An important process can be given more lottery tickets
 - Can control the proportion of usage
- ❑ Each resource can have its own set of tickets
 - Different proportion of usage **per resource per task**
- ❑ Simple implementation

Summary

■ Scheduling in OS:

- ❑ Basic definition
- ❑ Factors that affect scheduling
 - Process, Environment
- ❑ Criteria of good scheduling

■ Scheduling Algorithms:

- ❑ FCFS, SJF, SRT for Batch Processing System
- ❑ RR, Priority, Multi-Level Queues and Lottery scheduling for Interactive System

References

- Modern Operating System (4th Edition)
 - **Chapter 2.4**
- Operating System Concepts (9th Edition)
 - **Chapter 6**
- Operating Systems: Three Easy Pieces
 - <http://pages.cs.wisc.edu/~remzi/OSTEP>
 - **Chapters 7, 8, 9**
 - Advanced (optional): chapter 10